

Indexes Binning Outliers and Sampling

October 8, 2019

0.1 Environment

```
In [3]: import numpy as np
import pandas as pd
PREVIOUS_MAX_ROWS = pd.options.display.max_rows
pd.options.display.max_rows = 20
np.random.seed(12345)
import matplotlib.pyplot as plt
plt.rc('figure', figsize=(10, 6))
np.set_printoptions(precision=4, suppress=True)
```

0.1.1 Rename axis indexes

```
In [8]: data=pd.DataFrame(np.arange(12).reshape((3,4)),
                           index=['ant', 'bee', 'cat'],
                           columns=['one', 'two', 'three', 'four'])
```

```
In [9]: data
```

```
Out[9]:
```

	one	two	three	four
ant	0	1	2	3
bee	4	5	6	7
cat	8	9	10	11

```
In [11]: transform=lambda x: x[:4].upper() # Define a function that changes indexes into upper
data.index.map(transform)
```

```
Out[11]: Index(['ANT', 'BEE', 'CAT'], dtype='object')
```

```
In [13]: data.index=data.index.map(transform)
data
```

```
Out[13]:
```

	one	two	three	four
ANT	0	1	2	3
BEE	4	5	6	7
CAT	8	9	10	11

```
In [14]: data.rename(index=str.title, columns=str.upper)
```

```
Out[14]:
```

	ONE	TWO	THREE	FOUR
Ant	0	1	2	3
Bee	4	5	6	7
Cat	8	9	10	11

```
In [18]: data.rename(index={'ANT':'Ant-eater'},
                      columns={'three':'two-n-a-half'})
```

```
Out[18]:
```

	one	two	two-n-a-half	four
Ant-eater	0	1	2	3
BEE	4	5	6	7
CAT	8	9	10	11

```
In [19]: data.rename(index={'ANT':'Antelope'}, inplace=True)
data
```

```
Out[19]:
```

	one	two	three	four
Antelope	0	1	2	3
BEE	4	5	6	7
CAT	8	9	10	11

0.1.2 Discretization and Binning

```
In [37]: ages = [20, 22, 25, 27, 21, 23, 37, 31, 61, 45, 41, 32] # alternatively can use ages=
ages
```

```
Out[37]: [20, 22, 25, 27, 21, 23, 37, 31, 61, 45, 41, 32]
```

```
In [39]: bins=[18,25,35,60,100]
cats=pd.cut(ages,bins) #categories with open lower but closed upper limits
cats
```

```
Out[39]: [(18, 25], (18, 25], (18, 25], (25, 35], (18, 25], ..., (25, 35], (60, 100], (35, 60]
Length: 12
Categories (4, interval[int64]): [(18, 25] < (25, 35] < (35, 60] < (60, 100]]
```

```
In [40]: cats.codes
```

```
Out[40]: array([0, 0, 0, 1, 0, 0, 2, 1, 3, 2, 2, 1], dtype=int8)
```

```
In [41]: cats.categories
```

```
Out[41]: IntervalIndex([(18, 25], (25, 35], (35, 60], (60, 100]],
                        closed='right',
                        dtype='interval[int64]')
```

```
In [42]: pd.value_counts(cats) # frequency counts
```

```
Out[42]: (18, 25]      5
(35, 60]      3
(25, 35]      3
(60, 100]      1
dtype: int64
```

```

In [43]: pd.cut(ages,[18,26,36,61,100],right=False) #categories with closed lower but open upper

Out[43]: [[18, 26), [18, 26), [18, 26), [26, 36), [18, 26), ..., [26, 36), [61, 100), [36, 61)]
Length: 12
Categories (4, interval[int64]): [[18, 26) < [26, 36) < [36, 61) < [61, 100)]

In [44]: group_names=['Youth', 'YoungAdult', 'MiddleAged', 'Senior']
pd.cut(ages, bins, labels=group_names)

Out[44]: [Youth, Youth, Youth, YoungAdult, Youth, ..., YoungAdult, Senior, MiddleAged, MiddleAged]
Length: 12
Categories (4, object): [Youth < YoungAdult < MiddleAged < Senior]

In [46]: data=np.random.rand(20) #uniform random 20 numbers
data

Out[46]: array([0.8374, 0.3832, 0.2988, 0.0063, 0.4376, 0.7379, 0.3758, 0.4932,
0.014 , 0.2494, 0.3471, 0.7967, 0.9384, 0.1005, 0.7354, 0.9764,
0.7059, 0.9518, 0.9278, 0.41  ])

In [48]: pd.cut(data, 4, precision=2)

Out[48]: [(0.73, 0.98], (0.25, 0.49], (0.25, 0.49], (0.0053, 0.25], (0.25, 0.49], ..., (0.73, 0.98]]
Length: 20
Categories (4, interval[float64]): [(0.0053, 0.25] < (0.25, 0.49] < (0.49, 0.73] < (0.73, 0.98]]

In [52]: data=np.random.randn(1000) #1000 random normal numbers
cats=pd.qcut(data,4) #Group into quartiles
cats
pd.value_counts(cats)

Out[52]: (0.663, 3.389]      250
(-0.0374, 0.663]      250
(-0.673, -0.0374]      250
(-3.665, -0.673]      250
dtype: int64

In [57]: newcats=pd.qcut(data,[0,0.2,0.5,0.8,1.])
newcats

Out[57]: [(-0.863, -0.0374], (-0.863, -0.0374], (-3.665, -0.863], (0.83, 3.389], (0.83, 3.389]]
Length: 1000
Categories (4, interval[float64]): [(-3.665, -0.863] < (-0.863, -0.0374] < (-0.0374, 0.83] < (0.83, 3.389]]

In [58]: pd.value_counts(newcats)

Out[58]: (-0.0374, 0.83]      300
(-0.863, -0.0374]      300
(0.83, 3.389]      200
(-3.665, -0.863]      200
dtype: int64

```

0.1.3 Identify and work with outliers

```
In [60]: data=pd.DataFrame(np.random.randn(1000, 4)) #data matrix with thousand rows and 4 columns  
data.describe() #summary statistics
```

```
Out [60]:
```

	0	1	2	3
count	1000.000000	1000.000000	1000.000000	1000.000000
mean	-0.060144	-0.009583	0.004146	-0.020177
std	1.025283	0.991398	1.010832	1.009819
min	-3.067430	-3.423739	-2.944923	-3.094371
25%	-0.756485	-0.671014	-0.674305	-0.688851
50%	-0.056539	0.032223	-0.015343	-0.026830
75%	0.651537	0.661970	0.673732	0.672980
max	3.571767	3.424722	3.893606	4.104784

```
In [62]: col=data[2]  
col[np.abs(col)>3] #find values greater than 3 in column 2
```

```
Out [62]: 447    3.354485  
583    3.893606  
Name: 2, dtype: float64
```

```
In [65]: data[(np.abs(data)>3).any(1)] #report rows that have >3 in any column
```

```
Out [65]:
```

	0	1	2	3
222	0.429820	-0.247168	-1.145995	4.104784
380	-3.067430	0.043376	0.709777	-1.326205
412	-0.008728	-3.423739	1.061722	-0.398055
447	-1.975929	1.117683	3.354485	-1.824912
460	0.337453	-1.199839	-0.140934	3.216015
561	3.571767	-0.080974	-0.362215	-1.887861
571	1.186184	3.162137	-1.811221	-0.295279
583	-0.477607	0.101242	3.893606	1.048426
629	0.111325	-0.379214	0.862023	-3.094371
977	-3.019376	-0.534652	1.155369	1.047623
986	0.189540	3.424722	0.871550	0.452319

```
In [73]: data[np.abs(data)>3]=np.sign(data)*3 #replace outliers with 3 (keep the sign)  
data
```

```
Out [73]:
```

	0	1	2	3
0	0.405399	2.695794	0.564636	1.593455
1	0.885846	0.324446	0.606096	0.915896
2	-1.030575	-1.402759	-0.910587	-0.956956
3	0.270255	-2.908266	0.460448	-2.941183
4	1.430784	0.694300	-0.236944	-0.588769
5	0.675679	-0.896593	-0.928476	-0.582420
6	-0.465101	-0.484181	-2.008272	1.902356
7	-0.075422	-1.194952	-0.955007	0.748794

```

8      0.434854 -1.295906 -0.060320  0.763275
9     -1.033090  0.277089 -0.866665  0.682160
..      ...      ...      ...      ...
990    0.243201  0.829935  0.662501  0.148890
991    1.358860 -1.216059 -0.703997 -2.141126
992    0.505308  0.188781 -0.541365 -0.686123
993    0.283435  0.345885  0.576676  0.388423
994   -0.004547 -0.626815 -1.372446  0.297942
995   -0.185451  0.130191 -0.424046  1.947049
996   -0.200710  0.909408 -0.599311 -1.033973
997   -0.538521 -2.125265 -0.475660 -0.774994
998    1.202417  0.082925 -0.858229 -1.828928
999   -0.327564 -0.267989  0.845160  0.371163

```

```
[1000 rows x 4 columns]
```

```
In [74]: np.sign(data).head() #check the sign only
```

```
Out[74]:
```

	0	1	2	3
0	1.0	1.0	1.0	1.0
1	1.0	1.0	1.0	1.0
2	-1.0	-1.0	-1.0	-1.0
3	1.0	-1.0	1.0	-1.0
4	1.0	1.0	-1.0	-1.0

0.1.4 Permutation and random sampling from the data

```
In [80]: df=pd.DataFrame(np.arange(5*4).reshape(5,4)) #create 5*4=20 consecutive values then r
df
```

```
Out[80]:
```

	0	1	2	3
0	0	1	2	3
1	4	5	6	7
2	8	9	10	11
3	12	13	14	15
4	16	17	18	19

```
In [86]: sampler=np.random.permutation(5) #random arrangements of rows
sampler
```

```
Out[86]: array([4, 1, 0, 3, 2])
```

```
In [87]: df.take(sampler) #arrange rows according to sampler
```

```
Out[87]:
```

	0	1	2	3
4	16	17	18	19
1	4	5	6	7
0	0	1	2	3
3	12	13	14	15
2	8	9	10	11

```
In [89]: df.sample(n=3) #now sample the first three rows
```

```
Out[89]:
```

0	1	2	3
1	4	5	6
2	8	9	10
0	0	1	2

```
In [93]: # Another example: randomly choose values with replacement  
choices=pd.Series([5,7,-1,6,4]) #some random values  
choices
```

```
Out[93]:
```

0	5
1	7
2	-1
3	6
4	4

dtype: int64

```
In [94]: draws=choices.sample(n=10, replace=True)  
draws
```

```
Out[94]:
```

0	5
1	7
0	5
4	4
4	4
3	6
4	4
4	4
4	4
4	4

dtype: int64